



Modeling with the `tidymodels` Framework

Justin Post

Modeling Process

Given a model, we **fit** the model using data

- Must determine how well the model predicts on **new** data
- Create a test set or use CV (or perhaps both...)
- Judge effectiveness using a **metric** on predictions made from the model

Preparing the Data

General flow for modeling

- Read data in
- EDA (or perhaps after train/test split...)
- Split data into train and test (do response transform first!)
- Modify training data set predictors as needed
 - Center/scale
 - Create factors & dummy variables
 - Create interactions/quadratics/etc.
 - Log transform
 - ...
- Fit model(s) on training data
- Use same transformations on the test data or in CV process (*exactly* as done in training set)
- Predict on the test set

Convert Data

- We saw the use of `rsample::initial_split()`
 - **If doing a *non-learned* transformation, do those first outside of `tidymodels`**

```
library(tidyverse)
library(tidymodels)
bike_data <- read_csv("https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv") |>
  mutate(log_selling_price = log(selling_price)) |>
  select(-selling_price)
#save creation of new variables for now!
bike_split <- initial_split(bike_data, prop = 0.7)
bike_train <- training(bike_split)
bike_test <- testing(bike_split)
```

- `initial_split()` allows for stratified sampling too!

Data Preparation with `tidymodels`

- `recipes` package within `tidymodels` allows for transformations
 - Process keeps track of proper values to use for you!
 - Start with ``recipe()`` call
 - Denote formula for response/predictors and data to use
 - `summary()` describes current setup (we don't want all of these as predictors)

```
recipe(log_selling_price ~ ., data = bike_train) |>
  summary()
```

```
## # A tibble: 7 × 4
##   variable      type      role      source
##   <chr>      <list>   <chr>    <chr>
## 1 name      <chr [3]> predictor original
## 2 year      <chr [2]> predictor original
## 3 seller_type <chr [3]> predictor original
## 4 owner      <chr [3]> predictor original
## 5 km_driven  <chr [2]> predictor original
## 6 ex_showroom_price <chr [2]> predictor original
## 7 log_selling_price <chr [2]> outcome  original
```

Data Prepration with `tidymodels`

- `recipes` package within `tidymodels` allows for transformations
 - `update_role()` allows you to declare types of variables (such as `ID`)
 - This keeps the variable around even when not used in a model

```
recipe(log_selling_price ~ ., data = bike_train) |>  
  update_role(name, new_role = "ID") |>  
  summary()
```

```
## # A tibble: 7 × 4  
##   variable      type      role      source  
##   <chr>      <list>   <chr>    <chr>  
## 1 name      <chr [3]> ID       original  
## 2 year      <chr [2]> predictor original  
## 3 seller_type <chr [3]> predictor original  
## 4 owner      <chr [3]> predictor original  
## 5 km_driven  <chr [2]> predictor original  
## 6 ex_showroom_price <chr [2]> predictor original  
## 7 log_selling_price <chr [2]> outcome  original
```

Now Add Transformation Steps

- Many `step_*` functions to consider
 - `step_log()` to create our `log_km_driven` variable
 - `step_rm()` to remove a variable
 - `step_dummy()` to create dummy values for categorical variables
 - `step_normalize()` to center and scale numeric predictors

```
recipe(log_selling_price ~ ., data = bike_train) |>  
  update_role(name, new_role = "ID") |>  
  step_log(km_driven) |>  
  step_rm(ex_showroom_price) |>#too many nas  
  step_dummy(owner, seller_type) |>  
  step_normalize(all_numeric(), -all_outcomes())
```

prep() & bake() the Recipe

- If you have at least one preprocessing operation, `prep()` 'estimates the required parameters from a training set that can be later applied to other data sets'
- `bake()` applies the computations to data

```
recipe(log_selling_price ~ ., data = bike_train) |>  
  update_role(name, new_role = "ID") |>  
  step_log(km_driven) |>  
  step_rm(ex_showroom_price) |>  
  step_dummy(owner, seller_type) |>  
  step_normalize(all_numeric(), -all_outcomes()) |>  
  prep(training = bike_train) |>  
  bake(bike_train)
```

```
## # A tibble: 742 × 8  
##   name      year km_driven log_selling_price owner_X2nd.owner owner_X3rd.owner  
##   <fct>    <dbl>    <dbl>         <dbl>         <dbl>         <dbl>  
## 1 Bajaj Di... -0.405     0.808          10.3         -0.367        -0.111  
## 2 Honda Ac...  0.274    -1.05           10.6         -0.367        -0.111  
## 3 Bajaj Pu... -1.99     -0.0115         9.80         -0.367        -0.111  
## 4 Hero HF ...  0.726     0.197          10.5         -0.367        -0.111  
## 5 Royal En... -0.179     0.788          11.4         -0.367        -0.111  
## # i 737 more rows  
## # i 2 more variables: owner_X4th.owner <dbl>, seller_type_Individual <dbl>
```


`parsnip` for Creating a Model

- `prep()` and `bake()` steps are not required but help us debug/see what things look like
- Once we have our `recipe()` ready, we also need to do our modeling setup
 - Use `parsnip` package to specify a model
 - `parsnip` abstracts away the individual package syntax
 - Specify the model type and model engine
 - **This page** allows us to search for a model type so we can see which `model` and `engine` we want to specify!

Creating a Model with `tidymodels`

- Fit MLR model with `linear_reg()`
- Engine set to `lm` for basic models
- [Info page](#)

```
linear_reg() %>%  
  set_engine("lm") %>%  
  translate()
```

```
## Linear Regression Model Specification (regression)
```

```
##
```

```
## Computational engine: lm
```

```
##
```

```
## Model fit template:
```

```
## stats::lm(formula = missing_arg(), data = missing_arg(), weights = missing_arg())
```

Creating a Model with `tidymodels`

- Set up our model and recipes

```
bike_rec <- recipe(log_selling_price ~ ., data = bike_train) |>  
  update_role(name, new_role = "ID") |>  
  step_log(km_driven) |>  
  step_rm(ex_showroom_price) |>  
  step_dummy(owner, seller_type) |>  
  step_normalize(all_numeric(), -all_outcomes())  
  
bike_mod <- linear_reg() %>%  
  set_engine("lm")
```

workflow()s with tidymodels

- Now we can create a `workflow()`
 - Add our recipe and model with their corresponding functions

```
bike_wfl <- workflow() |>  
  add_recipe(bike_rec) |>  
  add_model(bike_mod)  
bike_wfl
```

```
## == Workflow ==  
## Preprocessor: Recipe  
## Model: linear_reg()  
##  
## — Preprocessor —  
## 4 Recipe Steps  
##  
## • step_log()  
## • step_rm()  
## • step_dummy()  
## • step_normalize()  
##  
## — Model —  
## Linear Regression Model Specification (regression)  
##  
## Computational engine: lm
```

fit() That Model!

- Finally, `fit()` allows us to fit our model to a data set!
- `tidy()` puts the results into a `tibble`

```
bike_fit <- bike_wfl |>  
  fit(bike_train)  
bike_fit |>  
  tidy()
```

```
## # A tibble: 7 × 5  
##   term                estimate std.error statistic  p.value  
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept)        10.7      0.0179     598.      0  
## 2 year                0.336      0.0210     16.0  1.00e-49  
## 3 km_driven          -0.241      0.0204    -11.8  1.80e-29  
## 4 owner_X2nd.owner    0.00538    0.0183     0.293  7.69e- 1  
## 5 owner_X3rd.owner    0.0740     0.0183     4.03  6.08e- 5  
## 6 owner_X4th.owner    0.0333     0.0180     1.85  6.52e- 2  
## 7 seller_type_Individual 0.00835    0.0180     0.464  6.43e- 1
```

Find Test Set Metric(s)

- Here we don't have a bunch of models we are comparing, only one is fit
- Can use `last_fit()` on the original `initial_split()` object to see how it performs
- `collect_metrics()` returns the metrics on the test set!

```
bike_fit |>
  last_fit(bike_split) |>
  collect_metrics()
```

```
## # A tibble: 2 × 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      0.556 Preprocessor1_Model1
## 2 rsq     standard      0.466 Preprocessor1_Model1
```

Find Test Set Metric(s)

- Here we don't have a bunch of models we are comparing, only one is fit
- Can use `last_fit()` on the original `initial_split()` object to see how it performs
- `collect_metrics()` returns the metrics on the test set!

```
bike_fit |>
  last_fit(bike_split) |>
  collect_metrics()
```

```
## # A tibble: 2 × 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      0.556 Preprocessor1_Model1
## 2 rsq     standard      0.466 Preprocessor1_Model1
```

The same transformations from the training set are used on the test set!

Fitting the Model with Cross-Validation

- Let's use 10 fold CV in the training set instead
 - Compare to another model's CV fit on the training set
 - Send best model to test set

Fitting the Model with Cross-Validation

- Let's use 10 fold CV in the training set instead
 - Compare to another model's CV fit on the training set
 - Send best model to test set
- Use `vfold_cv()` to split the data up and use `fit_resamples()` to fit the model appropriately

```
bike_10_fold <- vfold_cv(bike_train, 10)
bike_CV_fits <- bike_wfl |>
  fit_resamples(bike_10_fold)
bike_CV_fits
```

```
## # Resampling results
## # 10-fold cross-validation
## # A tibble: 10 × 4
##   splits          id    .metrics      .notes
##   <list>         <chr> <list>      <list>
## 1 <split [667/75]> Fold01 <tibble [2 × 4]> <tibble [0 × 3]>
## 2 <split [667/75]> Fold02 <tibble [2 × 4]> <tibble [0 × 3]>
## 3 <split [668/74]> Fold03 <tibble [2 × 4]> <tibble [0 × 3]>
## 4 <split [668/74]> Fold04 <tibble [2 × 4]> <tibble [0 × 3]>
## 5 <split [668/74]> Fold05 <tibble [2 × 4]> <tibble [0 × 3]>
## 6 <split [668/74]> Fold06 <tibble [2 × 4]> <tibble [0 × 3]>
## 7 <split [668/74]> Fold07 <tibble [2 × 4]> <tibble [0 × 3]>
## 8 <split [668/74]> Fold08 <tibble [2 × 4]> <tibble [0 × 3]>
```

Fitting the Model with Cross-Validation

- Combine the metrics using `collect_metrics()`

```
bike_CV_fits |>  
  collect_metrics()
```

```
## # A tibble: 2 × 6  
##   .metric .estimator  mean     n std_err .config  
##   <chr>   <chr>      <dbl> <int>  <dbl> <chr>  
## 1 rmse    standard    0.495     10  0.0257 Preprocessor1_Model1  
## 2 rsq     standard    0.498     10  0.0462 Preprocessor1_Model1
```

- This is our CV error on the training set!

Fit another Model with Cross-Validation for Comparison

- Let's build another recipe that includes interaction terms

```
bike_int_rec <- recipe(log_selling_price ~ ., data = bike_train) |>  
  update_role(name, new_role = "ID") |>  
  step_log(km_driven) |>  
  step_rm(ex_showroom_price) |>  
  step_dummy(owner, seller_type) |>  
  step_normalize(all_numeric(), -all_outcomes()) |>  
  step_interact(terms = ~km_driven*year*starts_with("seller_type"))
```

Fit another Model with Cross-Validation for Comparison

- Fit the model to the resamples and see our metric

```
bike_int_CV_fits <- workflow() |>
  add_recipe(bike_int_rec) |>
  add_model(bike_mod) |>
  fit_resamples(bike_10_fold)
rbind(bike_CV_fits |> collect_metrics(),
      bike_int_CV_fits |> collect_metrics())
```

```
## # A tibble: 4 × 6
##   .metric .estimator mean      n std_err .config
##   <chr>   <chr>      <dbl> <int>   <dbl> <chr>
## 1 rmse    standard    0.495     10  0.0257 Preprocessor1_Model1
## 2 rsq      standard    0.498     10  0.0462 Preprocessor1_Model1
## 3 rmse    standard    0.540     10  0.0390 Preprocessor1_Model1
## 4 rsq      standard    0.460     10  0.0505 Preprocessor1_Model1
```

- Simpler model is better here
- Could now compare its performance on the test set to some other 'best' models

Recap

- `tidymodels` provides a framework for predictive modeling
- Define a recipe
- Define a model and engine
- Fit the models (perhaps using cross-validation) and investigate metrics