

HW 8 & 9 -Modeling Practice

This homework is meant to give you a chance to do some structured practice with fitting linear models in R.

Data

We will use a dataset from the UCI Machine Learning Repository. This data set is about bike sharing rentals and is available at the assignment link. You can learn more about the data [here](https://www4.stat.ncsu.edu/~online/datasets/SeoulBikeData.csv). The data is available at <https://www4.stat.ncsu.edu/~online/datasets/SeoulBikeData.csv>

The data description describes the following variables:

- Date : **day/month/year**
- Rented Bike count - Count of bikes rented at each hour
- Hour - Hour of the day
- Temperature-Temperature in Celsius
- Humidity - %
- Windspeed - m/s
- Visibility - 10m
- Dew point temperature - Celsius
- Solar radiation - MJ/m2
- Rainfall - mm
- Snowfall - cm
- Seasons - Winter, Spring, Summer, Autumn
- Holiday - Holiday/No holiday
- Functional Day - NoFunc(Non Functional Hours), Fun(Functional hours)

Reading Data

- First read in the data
- When using `readr::read_csv()` I got an error `Error in nchar(x, "width") : invalid multibyte string, element 1`
- Google this and it is a quick fix!

```
library(tidyverse)
library(tidymodels)
bike_data <- read_csv("https://www4.stat.ncsu.edu/~online/datasets/SeoulBikeData.csv",
                      local = locale(encoding = "latin1"))
bike_data
```

```
## # A tibble: 8,760 x 14
##   Date      'Rented Bike Count' Hour 'Temperature(°C)' 'Humidity(%)'
##   <chr>          <dbl> <dbl>          <dbl>          <dbl>
## 1 01/12/2017      254     0          -5.2           37
## 2 01/12/2017      204     1          -5.5           38
```

```
## 3 01/12/2017      173      2      -6      39
## 4 01/12/2017      107      3     -6.2     40
## 5 01/12/2017       78      4      -6      36
## 6 01/12/2017      100      5     -6.4      37
## 7 01/12/2017      181      6     -6.6      35
## 8 01/12/2017      460      7     -7.4      38
## 9 01/12/2017      930      8     -7.6      37
## 10 01/12/2017     490      9     -6.5      27
## # i 8,750 more rows
## # i 9 more variables: 'Wind speed (m/s)' <dbl>, 'Visibility (10m)' <dbl>,
## #   'Dew point temperature(°C)' <dbl>, 'Solar Radiation (MJ/m2)' <dbl>,
## #   'Rainfall(mm)' <dbl>, 'Snowfall (cm)' <dbl>, Seasons <chr>, Holiday <chr>,
## #   'Functioning Day' <chr>
```

EDA

Checking the Data

First, let's check for missingness (there's a `tidy` way to do this but this is just a `baseR` way).

```
bike_data |>
  is.na() |>
  colSums()
```

```
##           Date      Rented Bike Count      Hour
##           0           0           0
## Temperature(°C)      Humidity(%)      Wind speed (m/s)
##           0           0           0
## Visibility (10m) Dew point temperature(°C)      Solar Radiation (MJ/m2)
##           0           0           0
## Rainfall(mm)      Snowfall (cm)      Seasons
##           0           0           0
## Holiday      Functioning Day
##           0           0
```

No apparent missingness. Let's check for column type and values.

```
attributes(bike_data)$spec
```

```
## cols(
##   Date = col_character(),
##   'Rented Bike Count' = col_double(),
##   Hour = col_double(),
##   'Temperature(°C)' = col_double(),
##   'Humidity(%)' = col_double(),
##   'Wind speed (m/s)' = col_double(),
##   'Visibility (10m)' = col_double(),
##   'Dew point temperature(°C)' = col_double(),
##   'Solar Radiation (MJ/m2)' = col_double(),
##   'Rainfall(mm)' = col_double(),
##   'Snowfall (cm)' = col_double(),
##   Seasons = col_character(),
```

```
## Holiday = col_character(),
## 'Functioning Day' = col_character()
## )
```

All columns seem reasonable except the `Date` column. Let's turn that into a real date.

```
bike_data <- bike_data |>
  mutate(date = lubridate::dmy(Date)) |>
  select(-Date)
```

Now briefly summarize each column to see if there are any weird values.

```
summary(bike_data)
```

```
## Rented Bike Count      Hour      Temperature(°C)  Humidity(%)
## Min.   : 0.0      Min.   : 0.00      Min.   : -17.80      Min.   : 0.00
## 1st Qu.: 191.0    1st Qu.: 5.75      1st Qu.: 3.50      1st Qu.: 42.00
## Median : 504.5    Median : 11.50     Median : 13.70     Median : 57.00
## Mean   : 704.6    Mean   : 11.50     Mean   : 12.88     Mean   : 58.23
## 3rd Qu.: 1065.2   3rd Qu.: 17.25     3rd Qu.: 22.50     3rd Qu.: 74.00
## Max.   : 3556.0   Max.   : 23.00     Max.   : 39.40     Max.   : 98.00
## Wind speed (m/s)  Visibility (10m)  Dew point temperature(°C)
## Min.   :0.000      Min.   : 27      Min.   : -30.600
## 1st Qu.:0.900      1st Qu.: 940      1st Qu.: -4.700
## Median :1.500      Median : 1698     Median : 5.100
## Mean   :1.725      Mean   : 1437     Mean   : 4.074
## 3rd Qu.:2.300      3rd Qu.: 2000     3rd Qu.: 14.800
## Max.   :7.400      Max.   : 2000     Max.   : 27.200
## Solar Radiation (MJ/m2)  Rainfall(mm)  Snowfall (cm)      Seasons
## Min.   :0.0000      Min.   : 0.0000      Min.   :0.00000      Length:8760
## 1st Qu.:0.0000      1st Qu.: 0.0000      1st Qu.:0.00000      Class :character
## Median :0.0100      Median : 0.0000      Median :0.00000      Mode  :character
## Mean   :0.5691      Mean   : 0.1487      Mean   :0.07507
## 3rd Qu.:0.9300      3rd Qu.: 0.0000      3rd Qu.:0.00000
## Max.   :3.5200      Max.   : 35.0000      Max.   :8.80000
## Holiday      Functioning Day      date
## Length:8760      Length:8760      Min.   :2017-12-01
## Class :character  Class :character  1st Qu.:2018-03-02
## Mode  :character  Mode  :character  Median :2018-06-01
##                                     Mean   :2018-06-01
##                                     3rd Qu.:2018-08-31
##                                     Max.   :2018-11-30
```

I don't know much about weather but things seem ok. Visibility (10m) is likely truncated at 2000.

Check the character columns

```
bike_data$Seasons |>
  unique()
```

```
## [1] "Winter" "Spring" "Summer" "Autumn"
```

```
bike_data$Holiday |>
  unique()
```

```
## [1] "No Holiday" "Holiday"
```

```
bike_data$`Functioning Day` |>
  unique()
```

```
## [1] "Yes" "No"
```

Ok, no worries there! Let's turn these into factor variables.

```
bike_data <- bike_data |>
  mutate(seasons = factor(Seasons),
         holiday = factor(Holiday),
         fn_day = factor(`Functioning Day`)) |>
  select(-Seasons, -Holiday, -`Functioning Day`)
```

Lastly, I think renaming the rest of the variables will be beneficial.

```
bike_data <- bike_data |>
  rename('bike_count' = `Rented Bike Count`,
        'hour' = "Hour",
        "temp" = `Temperature(°C)`,
        "wind_speed" = `Wind speed (m/s)`,
        "humidity" = `Humidity(%)`,
        "vis" = `Visibility (10m)`,
        "dew_point_temp" = `Dew point temperature(°C)`,
        "solar_radiation" = `Solar Radiation (MJ/m2)`,
        "rainfall" = "Rainfall(mm)",
        "snowfall" = `Snowfall (cm)`)
```

- Seems like the `fn_day` variable implies they were out of commission sometimes. Let's remove those observations and that variable.

```
bike_data <- bike_data |>
  filter(fn_day == "Yes") |>
  select(-fn_day)
```

To simplify our analysis, we'll summarize across the hours so that each day has one observation associated with it. Let's `group_by()` the `date`, `seasons`, and `holiday` variables and find the sum of the `bike_count`, `rainfall`, and `snowfall` variables and the mean of all the weather related variables.

```
bike_data <- bike_data |>
  group_by(date, seasons, holiday) |>
  summarize(bike_count = sum(bike_count),
           temp = mean(temp),
           humidity = mean(humidity),
           wind_speed = mean(wind_speed),
           vis = mean(vis),
```

```

    dew_point_temp = mean(dew_point_temp),
    solar_radiation = mean(solar_radiation),
    rainfall = sum(rainfall),
    snowfall = sum(snowfall)) |>
ungroup()

```

```
bike_data
```

```

## # A tibble: 353 x 12
##   date      seasons holiday bike_count temp humidity wind_speed vis
##   <date>    <fct>   <fct>    <dbl>  <dbl>   <dbl>   <dbl> <dbl>
## 1 2017-12-01 Winter No Holiday    9539 -2.45    45.9    1.54  1871.
## 2 2017-12-02 Winter No Holiday    8523  1.32    62.0    1.71  1471.
## 3 2017-12-03 Winter No Holiday    7222  4.88    81.5    1.61   456.
## 4 2017-12-04 Winter No Holiday    8729 -0.304   52.5    3.45  1363.
## 5 2017-12-05 Winter No Holiday    8307 -4.46    36.4    1.11  1959.
## 6 2017-12-06 Winter No Holiday    6669  0.0458  70.8    0.696 1187.
## 7 2017-12-07 Winter No Holiday    8549  1.09    67.5    1.69   949.
## 8 2017-12-08 Winter No Holiday    8032 -3.82    41.8    1.85  1872.
## 9 2017-12-09 Winter No Holiday    7233 -0.846   46     1.08  1861.
## 10 2017-12-10 Winter No Holiday    3453  1.19    69.7    2.00  1043.
## # i 343 more rows
## # i 4 more variables: dew_point_temp <dbl>, solar_radiation <dbl>,
## #   rainfall <dbl>, snowfall <dbl>

```

Summary Stats & Graphs

Some quick summary stats. We're going to focus on modeling the `bike_count` so let's focus there.

Numeric summaries first. Let's produce the mean, median, sd, IQR, min, and max for this variable. Then do the same across levels of the categorical variables.

```

bike_data |>
  summarize(across(`bike_count`,
    .fns = c("mean" = mean,
             "median" = median,
             "sd" = sd,
             "IQR" = IQR,
             "min" = min,
             "max" = max),
    .names = "{.col}_{.fn}"))

## # A tibble: 1 x 6
##   bike_count_mean bike_count_median bike_count_sd bike_count_IQR bike_count_min
##   <dbl>          <dbl>          <dbl>          <dbl>          <dbl>
## 1    17485.        18563          9937.         19318          977
## # i 1 more variable: bike_count_max <dbl>

```

- Looks to be right skewed with a pretty large standard deviation.

```
bike_data |>
  group_by(holiday) |>
  summarize(across(`bike_count`,
    .fns = c("mean" = mean,
              "median" = median,
              "sd" = sd,
              "IQR" = IQR,
              "min" = min,
              "max" = max),
    .names = "{.col}_{.fn}"))
```

```
## # A tibble: 2 x 7
##   holiday    bike_count_mean bike_count_median bike_count_sd bike_count_IQR
##   <fct>          <dbl>          <dbl>          <dbl>          <dbl>
## 1 Holiday          12700.            7184          10504.          16576
## 2 No Holiday        17727.           19104.           9862.          19168.
## # i 2 more variables: bike_count_min <dbl>, bike_count_max <dbl>
```

```
bike_data |>
  group_by(seasons) |>
  summarize(across(`bike_count`,
    .fns = c("mean" = mean,
              "median" = median,
              "sd" = sd,
              "IQR" = IQR,
              "min" = min,
              "max" = max),
    .names = "{.col}_{.fn}"))
```

```
## # A tibble: 4 x 7
##   seasons    bike_count_mean bike_count_median bike_count_sd bike_count_IQR
##   <fct>          <dbl>          <dbl>          <dbl>          <dbl>
## 1 Autumn          22099.           23350           6711.          10733
## 2 Spring          17910.           17590           8357.          14362.
## 3 Summer          24818.           25572.           7297.           9308.
## 4 Winter           5413.            5498           1808.           2634.
## # i 2 more variables: bike_count_min <dbl>, bike_count_max <dbl>
```

- Strong differences depending on holiday and seasons.

```
bike_data |>
  group_by(seasons, holiday) |>
  summarize(across(`bike_count`,
    .fns = c("mean" = mean,
              "median" = median,
              "sd" = sd,
              "IQR" = IQR,
              "min" = min,
              "max" = max),
    .names = "{.col}_{.fn}"))
```

```
## # A tibble: 8 x 8
```

```
## # Groups:   seasons [4]
##   seasons holiday bike_count_mean bike_count_median bike_count_sd bike_count_IQR
##   <fct>   <fct>         <dbl>             <dbl>         <dbl>         <dbl>
## 1 Autumn  Holiday      22754.             21705         5642.         5740
## 2 Autumn  No Hol~      22065.             23472         6792.        10734
## 3 Spring  Holiday      15247.             13790        10917.        10844
## 4 Spring  No Hol~      18002.             17730         8322.        14224.
## 5 Summer  Holiday      24532.             24532.         8438.         5966.
## 6 Summer  No Hol~      24824.             25572.         7324.         9165
## 7 Winter  Holiday       3759              3454.         1561.         1060.
## 8 Winter  No Hol~       5574.             5609         1757.         2564
## # i 2 more variables: bike_count_min <dbl>, bike_count_max <dbl>
```

- Differences are pretty big in the Winter and Autumn but not the other seasons. Perhaps an interaction between these two variables is important.

Now let's do some correlation.


```
bike_data |>
  select(where(is.numeric)) |>
  cor() |>
  round(3)
```

```
##           bike_count  temp humidity wind_speed  vis dew_point_temp
## bike_count      1.000  0.753   0.036   -0.193  0.166         0.650
## temp            0.753  1.000   0.404   -0.261  0.002         0.963
## humidity        0.036  0.404   1.000   -0.234 -0.559         0.632
## wind_speed      -0.193 -0.261  -0.234    1.000  0.206        -0.288
## vis             0.166  0.002  -0.559    0.206  1.000        -0.154
## dew_point_temp   0.650  0.963   0.632   -0.288 -0.154         1.000
## solar_radiation  0.736  0.550  -0.274    0.096  0.271         0.383
## rainfall        -0.239  0.145   0.529   -0.102 -0.222         0.265
## snowfall        -0.265 -0.267   0.065    0.021 -0.102        -0.210
##
## solar_radiation rainfall snowfall
## bike_count      0.736  -0.239  -0.265
## temp            0.550   0.145  -0.267
## humidity        -0.274   0.529   0.065
## wind_speed      0.096  -0.102   0.021
## vis             0.271  -0.222  -0.102
## dew_point_temp   0.383   0.265  -0.210
## solar_radiation  1.000  -0.323  -0.233
## rainfall        -0.323   1.000  -0.023
## snowfall        -0.233  -0.023   1.000
```

- Definitely a few moderate relationships with bike_count here (temp and solar_radiation). temp and dew_point_temp are obviously pretty related. humidity and vis along with humidity and dew_point_temp as well.


Let's do some visualizations.

```
ggplot(bike_data, aes(x = temp, y = bike_count)) +
  geom_jitter(aes(color = seasons)) +
  facet_grid(~holiday)
```



Some expected trends here and we can see that once it gets pretty hot, bike rentals slow.

```
ggplot(bike_data, aes(x = solar_radiation, y = bike_count)) +  
  geom_point(aes(color = seasons)) +  
  facet_grid(~holiday)
```

More solar radiation is associated with more bike rentals. (Let's just remove the ozone!)

Ok, we could keep going with these types of plots to help us understand our data but I'm going to stop there.

Split the Data

- Use functions from `tidymodels` to split the data into a training and test set (75/25 split). Use the `strata` argument to stratify the split on the `seasons`.
- On the training set, create a 10 fold CV split

```
set.seed(11)
bike_split <- initial_split(bike_data, prop = 0.75, strata = seasons)
bike_train <- training(bike_split)
bike_test <- testing(bike_split)
bike_10_fold <- vfold_cv(bike_train, 10)
```

Fitting MLR Models

First, let's create some recipes.

For the 1st recipe:

- Let's ignore the `date` variable (so we'll need to remove that or give it a different ID) but use it to create a weekday/weekend (factor) variable. (See step 2 of the `shinyml` tutorial! You can use

`step_date()` then `step_mutate()` with an `factor(if_else(...))` to create the variable. I then had to remove the intermediate variable created.)

- Let's standardize the numeric variables since their scales are pretty different.
- Let's create dummy variables for the seasons, holiday, and our new day type variable

```
MLR_rec1 <- recipe(bike_count ~ ., data = bike_train) |>
  step_date(date, features = "dow") |>
  step_mutate(day_type = factor(if_else(date_dow %in% c("Sat", "Sun"), "Weekend", "Weekday"))) |>
  step_rm(date, date_dow) |>
  step_dummy(seasons, holiday, day_type) |>
  step_normalize(all_numeric(), -bike_count)
```

For the 2nd recipe:

- Do the same steps as above.
- Add in interactions between seasons and holiday, seasons and temp, temp and rainfall. For the seasons interactions, you can use `starts_with()` to create the proper interactions.

```
MLR_rec2 <- MLR_rec1 |>
  step_interact(terms = ~starts_with("seasons_")*starts_with("holiday_") +
                 starts_with("seasons_W")*temp +
                 temp*rainfall)
MLR_rec2 <- MLR_rec1 |>
  step_interact(terms = ~starts_with("seasons_")*starts_with("holiday_") + temp*rainfall) |>
  step_interact(terms = ~seasons_Spring*temp + seasons_Winter*temp) |>
  step_interact(terms = ~seasons_Summer*temp)
```

For the 3rd recipe:

- Do the same as the 2nd recipe.
- Add in quadratic terms for each numeric predictor

```
MLR_rec3 <- MLR_rec2 |>
  step_poly(temp,
            wind_speed,
            vis,
            dew_point_temp,
            solar_radiation,
            rainfall,
            snowfall,
            degree = 2, keep_original_cols = FALSE)
```

Now we can set up our linear model fit.

```
MLR_spec <- linear_reg() |>
  set_engine("lm")
```

Fit the models using 10 fold CV and consider the training set CV error to choose a best model.

```
MLR_CV_fit1 <- workflow() |>
  add_recipe(MLR_rec1) |>
  add_model(MLR_spec) |>
  fit_resamples(bike_10_fold)
MLR_CV_fit2 <- workflow() |>
  add_recipe(MLR_rec2) |>
  add_model(MLR_spec) |>
  fit_resamples(bike_10_fold)
MLR_CV_fit3 <- workflow() |>
  add_recipe(MLR_rec3) |>
  add_model(MLR_spec) |>
  fit_resamples(bike_10_fold)
```

Get our metrics:

```
rbind(MLR_CV_fit1 |> collect_metrics(),
      MLR_CV_fit2 |> collect_metrics(),
      MLR_CV_fit3 |> collect_metrics())
```

```
## # A tibble: 6 x 6
##   .metric .estimator    mean     n std_err .config
##   <chr>   <chr>         <dbl> <int>   <dbl> <chr>
## 1 rmse    standard    4284.     10  165.    Preprocessor1_Model11
## 2 rsq     standard     0.822     10  0.0151 Preprocessor1_Model11
## 3 rmse    standard    3156.     10  267.    Preprocessor1_Model11
## 4 rsq     standard     0.898     10  0.0176 Preprocessor1_Model11
## 5 rmse    standard    3070.     10  213.    Preprocessor1_Model11
## 6 rsq     standard     0.903     10  0.0142 Preprocessor1_Model11
```

The last model appears to be the best! Let's fit that to the entire training set and then see how it performs on the test set.

```
final_fit <- workflow() |>
  add_recipe(MLR_rec1) |>
  add_model(MLR_spec) |>
  last_fit(bike_split)
final_fit |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>         <dbl> <chr>
## 1 rmse    standard    3980.    Preprocessor1_Model11
## 2 rsq     standard     0.846    Preprocessor1_Model11
```

Obtain the final model (fit on the entire training set) coefficient table using `tidy()`.

```
final_fit |>
  extract_fit_parsnip() |>
  tidy()
```

```
## # A tibble: 14 x 5
##   term                estimate std.error statistic  p.value
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)      17446.      252.     69.3  9.38e-165
## 2 temp             -2439.     5215.    -0.468  6.40e- 1
## 3 humidity         -1927.     1904.    -1.01  3.13e- 1
## 4 wind_speed        -523.      286.    -1.83  6.86e- 2
## 5 vis               -63.7      361.    -0.177  8.60e- 1
## 6 dew_point_temp    7143.     6143.     1.16  2.46e- 1
## 7 solar_radiation   4088.      473.     8.64  6.74e-16
## 8 rainfall         -1779.      333.    -5.35  2.00e- 7
## 9 snowfall          -317.      276.    -1.15  2.50e- 1
##10 seasons_Spring  -2528.      355.    -7.12  1.14e-11
##11 seasons_Summer  -1670.      442.    -3.78  1.98e- 4
##12 seasons_Winter -3684.      501.    -7.35  2.88e-12
##13 holiday_No.Holiday  835.      256.     3.26  1.28e- 3
##14 day_type_Weekend -1050.      256.    -4.10  5.56e- 5
```

Fitting LASSO

We'll use the first recipe here.

Let's set up our model specification.

```
LASSO_spec <- linear_reg(penalty = tune(), mixture = 1) |>
  set_engine("glmnet")
```

Now we can create our workflow.


```
LASSO_wkf <- workflow() |>
  add_recipe(MLR_rec1) |>
  add_model(LASSO_spec)
```

Now we'll fit the model with `tune_grid()` and `grid_regular()`.

```
LASSO_grid <- LASSO_wkf |>
  tune_grid(resamples = bike_10_fold,
    grid = grid_regular(penalty(), levels = 400))
```

Let's see which tuning parameter is the best.

```
LASSO_grid |>
  collect_metrics() |>
  filter(.metric == "rmse") |>
  ggplot(aes(penalty, mean, color = .metric)) +
  geom_line()
```



Looks like all the penalty values give the same `rmse`. That means we'll select the smallest penalty.

```
lowest_rmse <- LASSO_grid |>
  select_best(metric = "rmse")
lowest_rmse
```

```
## # A tibble: 1 x 2
##   penalty .config
##   <dbl> <chr>
## 1 0.0000000001 Preprocessor1_Model001
```

Fitting Regression Tree

Let's set up our model specification.

```
tree_mod <- decision_tree(tree_depth = tune(),
                           min_n = 20,
                           cost_complexity = tune()) |>
  set_engine("rpart") |>
  set_mode("regression")
```

Now we can create our workflow.

```
tree_wkf <- workflow() |>
  add_recipe(MLR_rec1) |>
  add_model(tree_mod)
```

Now we'll fit the model with `tune_grid()` and `grid_regular()`.

```
tree_grid <- grid_regular(cost_complexity(),
                          tree_depth(),
                          levels = c(10, 5))
tree_fits <- tree_wkf |>
  tune_grid(resamples = bike_10_fold,
            grid = tree_grid)
```

Let's see which tuning parameter is the best.

```
tree_fits |>
  collect_metrics() |>
  filter(.metric == "rmse") |>
  arrange(mean)
```

```
## # A tibble: 50 x 8
##   cost_complexity tree_depth .metric .estimator mean    n std_err .config
##           <dbl>      <int> <chr>   <chr>      <dbl> <int>  <dbl> <chr>
## 1      0.001           11 rmse    standard  3817.    10   285. Preprocess~
## 2      0.001           15 rmse    standard  3817.    10   285. Preprocess~
## 3      0.0000000001     11 rmse    standard  3837.    10   302. Preprocess~
## 4      0.0000000001     11 rmse    standard  3837.    10   302. Preprocess~
## 5      0.000000001      11 rmse    standard  3837.    10   302. Preprocess~
## 6      0.0000001       11 rmse    standard  3837.    10   302. Preprocess~
## 7      0.000001        11 rmse    standard  3837.    10   302. Preprocess~
## 8      0.00001         11 rmse    standard  3837.    10   302. Preprocess~
## 9      0.0001          11 rmse    standard  3837.    10   302. Preprocess~
## 10     0.0000000001     15 rmse    standard  3837.    10   302. Preprocess~
## # i 40 more rows
```

Grab that tuning specification.

```
lowest_rmse_tree <- tree_fits |>
  select_best(metric = "rmse")
lowest_rmse_tree
```

```
## # A tibble: 1 x 3
##   cost_complexity tree_depth .config
##           <dbl>      <int> <chr>
## 1      0.001           11 Preprocessor1_Model38
```

Fitting Bagged Tree

Let's set up our model specification.

```
bag_spec <- bag_tree(tree_depth = 5, min_n = 10, cost_complexity = tune()) |>
  set_engine("rpart") |>
  set_mode("regression")
```

Now we can create our workflow.

```
library(baguette)
bag_wkf <- workflow() |>
  add_recipe(MLR_rec1) |>
  add_model(bag_spec)
```

Now we'll fit the model with `tune_grid()` and `grid_regular()`.

```
bag_fit <- bag_wkf |>
  tune_grid(resamples = bike_10_fold,
            grid = grid_regular(cost_complexity(),
                                levels = 15))
```

Let's see which tuning parameter is the best.

```
bag_fit |>
  collect_metrics() |>
  filter(.metric == "rmse") |>
  arrange(mean)
```

```
## # A tibble: 15 x 7
##   cost_complexity .metric .estimator  mean      n std_err .config
##   <dbl> <chr> <chr> <dbl> <int> <dbl> <chr>
## 1 1.64e- 7 rmse standard 3210.    10 168. Preprocessor1_Model06
## 2 1.39e- 5 rmse standard 3247.    10 194. Preprocessor1_Model09
## 3 8.48e- 9 rmse standard 3265.    10 173. Preprocessor1_Model04
## 4 7.20e- 7 rmse standard 3271.    10 179. Preprocessor1_Model07
## 5 6.11e- 5 rmse standard 3279.    10 182. Preprocessor1_Model10
## 6 3.16e- 6 rmse standard 3338.    10 185. Preprocessor1_Model08
## 7 4.39e-10 rmse standard 3341.    10 168. Preprocessor1_Model02
## 8 1.18e- 3 rmse standard 3353.    10 192. Preprocessor1_Model12
## 9 2.68e- 4 rmse standard 3364.    10 168. Preprocessor1_Model11
## 10 3.73e- 8 rmse standard 3371.    10 188. Preprocessor1_Model05
## 11 1.93e- 9 rmse standard 3413.    10 202. Preprocessor1_Model03
## 12 1 e-10 rmse standard 3471.    10 182. Preprocessor1_Model01
## 13 5.18e- 3 rmse standard 3489.    10 172. Preprocessor1_Model13
## 14 2.28e- 2 rmse standard 3980.    10 135. Preprocessor1_Model14
## 15 1 e- 1 rmse standard 4961.    10 173. Preprocessor1_Model15
```

Get the best one

```
lowest_rmse_bag <- bag_fit |>
  select_best(metric = "rmse")
lowest_rmse_bag
```

```
## # A tibble: 1 x 2
##   cost_complexity .config
##   <dbl> <chr>
## 1 0.000000164 Preprocessor1_Model06
```

Fitting Random Forest

```
rf_spec <- rand_forest(mtry = tune()) |>
  set_engine("ranger", importance = "impurity") |>
  set_mode("regression")
```

Now we can create our workflow.

```
rf_wkf <- workflow() |>
  add_recipe(MLR_rec1) |>
  add_model(rf_spec)
```

Now we'll fit the model with `tune_grid()` and `grid_regular()`.

```
rf_fit <- rf_wkf |>
  tune_grid(resamples = bike_10_fold,
            grid = 12)
```

Let's see which tuning parameter is the best.

```
rf_fit |>
  collect_metrics() |>
  filter(.metric == "rmse") |>
  arrange(mean)
```

```
## # A tibble: 9 x 7
##   mtry .metric .estimator mean      n std_err .config
##   <int> <chr>   <chr>    <dbl> <int>  <dbl> <chr>
## 1    13 rmse    standard  2979.    10   196. Preprocessor1_Model11
## 2    11 rmse    standard  3006.    10   200. Preprocessor1_Model12
## 3     9 rmse    standard  3019.    10   196. Preprocessor1_Model16
## 4    10 rmse    standard  3030.    10   192. Preprocessor1_Model15
## 5     7 rmse    standard  3045.    10   193. Preprocessor1_Model13
## 6     6 rmse    standard  3057.    10   196. Preprocessor1_Model19
## 7     4 rmse    standard  3145.    10   188. Preprocessor1_Model14
## 8     2 rmse    standard  3477.    10   171. Preprocessor1_Model18
## 9     1 rmse    standard  4784.    10   172. Preprocessor1_Model17
```

Get the best one

```
lowest_rmse_rf <- rf_fit |>
  select_best(metric = "rmse")
lowest_rmse_rf
```

```
## # A tibble: 1 x 2
##   mtry .config
##   <int> <chr>
## 1    13 Preprocessor1_Model11
```


Getting the mae for the Final Models and Other Summaries

Refit the MLR model and get all the rmse and mae.

```
final_fit <- workflow() |>
  add_recipe(MLR_rec1) |>
  add_model(MLR_spec) |>
  last_fit(bike_split, metrics = metric_set(rmse, mae))
final_fit |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>         <dbl> <chr>
## 1 rmse    standard         3980. Preprocessor1_Model11
## 2 mae     standard         3039. Preprocessor1_Model11
```

Let's get the coefficient table too.

```
workflow() |>
  add_recipe(MLR_rec1) |>
  add_model(MLR_spec) |>
  fit(bike_train) |>
  tidy()
```

```
## # A tibble: 14 x 5
##   term                estimate std.error statistic  p.value
##   <chr>                <dbl>   <dbl>    <dbl>   <dbl>
## 1 (Intercept)        17446.    252.    69.3  9.38e-165
## 2 temp               -2439.   5215.   -0.468 6.40e- 1
## 3 humidity           -1927.   1904.   -1.01  3.13e- 1
## 4 wind_speed          -523.    286.   -1.83  6.86e- 2
## 5 vis                 -63.7    361.   -0.177 8.60e- 1
## 6 dew_point_temp      7143.   6143.    1.16  2.46e- 1
## 7 solar_radiation     4088.    473.    8.64  6.74e-16
## 8 rainfall           -1779.    333.   -5.35  2.00e- 7
## 9 snowfall            -317.    276.   -1.15  2.50e- 1
##10 seasons_Spring     -2528.    355.   -7.12  1.14e-11
##11 seasons_Summer     -1670.    442.   -3.78  1.98e- 4
##12 seasons_Winter     -3684.    501.   -7.35  2.88e-12
##13 holiday_No.Holiday   835.    256.    3.26  1.28e- 3
##14 day_type_Weekend   -1050.    256.   -4.10  5.56e- 5
```

Let's fit our LASSO model to our training data and test on the test set, specifying we want rmse and mae.

```
LASSO_final <- LASSO_wkf |>
  finalize_workflow(lowest_rmse) |>
  last_fit(bike_split, metrics = metric_set(rmse, mae))
LASSO_final |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      3999. Preprocessor1_Model1
## 2 mae     standard      3063. Preprocessor1_Model1
```

Get the coefficient table too.

```
LASSO_wkf |>
  finalize_workflow(lowest_rmse) |>
  fit(bike_train) |>
  tidy()
```

```
## # A tibble: 14 x 3
##   term                estimate      penalty
##   <chr>              <dbl>      <dbl>
## 1 (Intercept)      17446. 0.0000000001
## 2 temp              389. 0.0000000001
## 3 humidity         -887. 0.0000000001
## 4 wind_speed       -522. 0.0000000001
## 5 vis               0 0.0000000001
## 6 dew_point_temp   3752. 0.0000000001
## 7 solar_radiation  4065. 0.0000000001
## 8 rainfall        -1841. 0.0000000001
## 9 snowfall         -336. 0.0000000001
## 10 seasons_Spring -2505. 0.0000000001
## 11 seasons_Summer -1607. 0.0000000001
## 12 seasons_Winter -3653. 0.0000000001
## 13 holiday_No.Holiday 820. 0.0000000001
## 14 day_type_Weekend -1060. 0.0000000001
```

Let's fit our regression tree model to our training data and test on the test set, specifying we want `rmse` and `mae`.

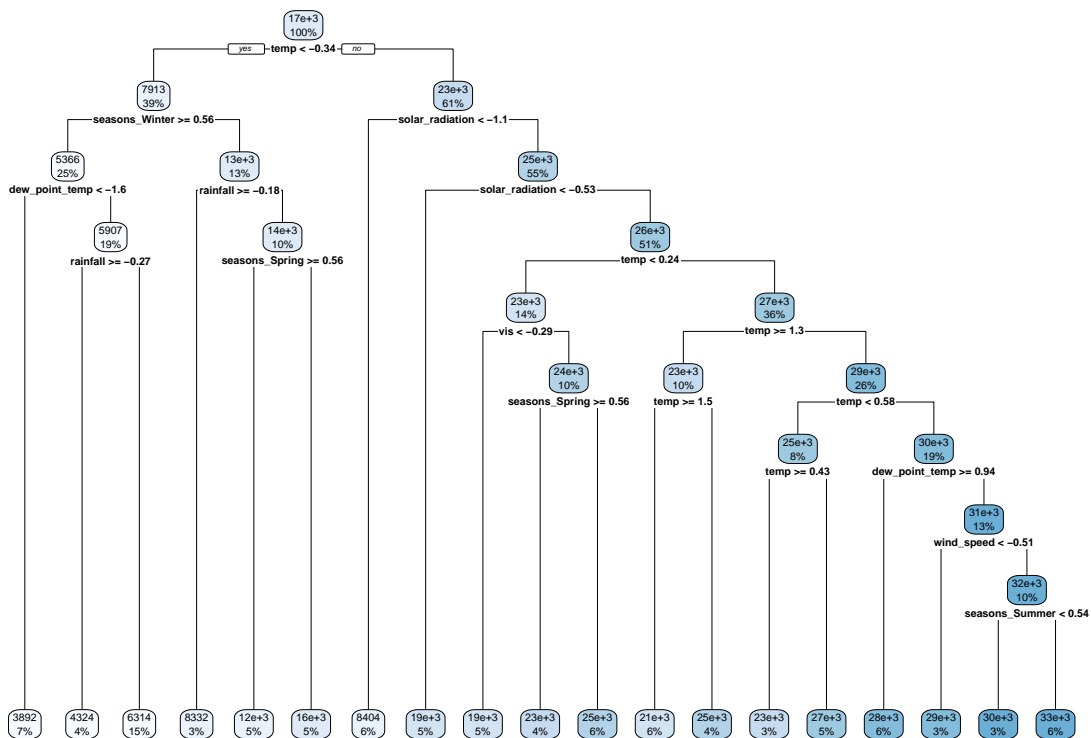
```
tree_final <- tree_wkf |>
  finalize_workflow(lowest_rmse_tree) |>
  last_fit(bike_split, metrics = metric_set(rmse, mae))
tree_final |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      3096. Preprocessor1_Model1
## 2 mae     standard      2362. Preprocessor1_Model1
```

Plot the final fit.

```
tree_final_train_fit <- extract_workflow(tree_final)

tree_final_train_fit %>%
  extract_fit_engine() %>%
  rpart.plot::rpart.plot(roundint = FALSE)
```




Let's fit our bagged tree model to our training data and test on the test set, specifying we want `rmse` and `mae`.

```
bag_final <- bag_wkf |>
  finalize_workflow(lowest_rmse_bag) |>
  last_fit(bike_split, metrics = metric_set(rmse, mae))
bag_final |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      3221. Preprocessor1_Model1
## 2 mae     standard      2516. Preprocessor1_Model1
```

Plot the variable importance.

```
bag_final_model <- extract_fit_engine(bag_final)
bag_final_model$imp |>
  mutate(term = factor(term, levels = term)) |>
  ggplot(aes(x = term, y = value)) +
  geom_bar(stat = "identity") +
  coord_flip()
```




Let's fit our rf model to our training data and test on the test set, specifying we want `rmse` and `mae`.

```
rf_final <- rf_wkf |>
  finalize_workflow(lowest_rmse_rf) |>
  last_fit(bike_split, metrics = metric_set(rmse, mae))
rf_final |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard         2621. Preprocessor1_Model11
## 2 mae     standard         2109. Preprocessor1_Model11
```

Plot the variable importance.

```
rf_final_model <- extract_fit_engine(rf_final)
tibble(term = names(rf_final_model$variable.importance),
       value = rf_final_model$variable.importance) |>
  arrange(value) |>
  mutate(term = factor(term, levels = term)) |>
  ggplot(aes(x = term, y = value)) +
  geom_bar(stat = "identity") +
  coord_flip()
```



Wrap up

Just to have all the test set metrics together:

```
LASSO_final |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      3999. Preprocessor1_Model11
## 2 mae     standard      3063. Preprocessor1_Model11
```

```
tree_final |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      3096. Preprocessor1_Model11
## 2 mae     standard      2362. Preprocessor1_Model11
```

```
bag_final |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      3221. Preprocessor1_Model1
## 2 mae     standard      2516. Preprocessor1_Model1
```

```
rf_final |>
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      2621. Preprocessor1_Model1
## 2 mae     standard      2109. Preprocessor1_Model1
```

Random forest is best on the test set. Let's fit it to the entire data set.

```
best_model <- rf_wkf |>
  finalize_workflow(lowest_rmse_rf) |>
  fit(bike_data)
```

```
best_model
```

```
## == Workflow [trained] =====
## Preprocessor: Recipe
## Model: rand_forest()
##
## -- Preprocessor -----
## 5 Recipe Steps
##
## * step_date()
## * step_mutate()
## * step_rm()
## * step_dummy()
## * step_normalize()
##
## -- Model -----
## Ranger result
##
## Call:
##   ranger::ranger(x = maybe_data_frame(x), y = y, mtry = min_cols(~13L,      x), importance = ~"impuri
##
## Type:                      Regression
## Number of trees:           500
## Sample size:               353
## Number of independent variables: 13
## Mtry:                       13
## Target node size:          5
```

## Variable importance mode:	impurity
## Splitrule:	variance
## OOB prediction error (MSE):	7591965
## R squared (OOB):	0.9231172