

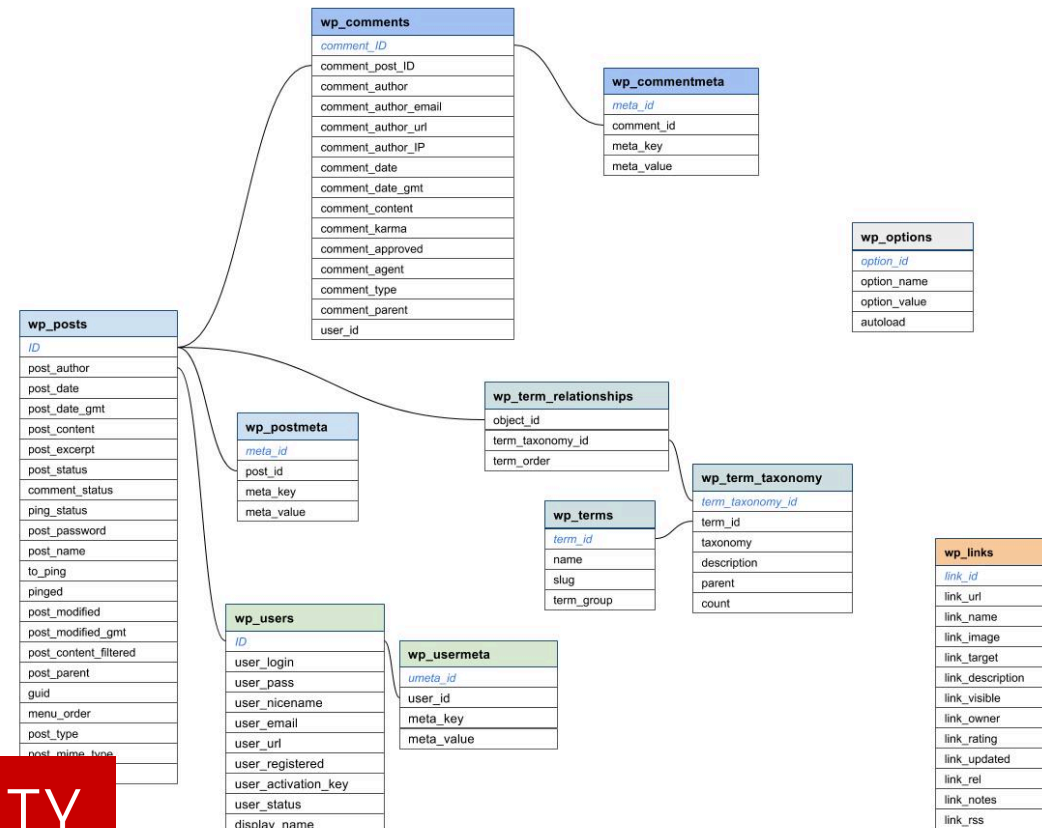


# SQL Style Joins

Justin Post

# Relational Databases

- Often want to combine data from multiple tables to summarize/model



- The common types of joins we do are given below! (Using `dplyr` not the particular SQL language.)

### Combine Data Sets

| x1 | x2 |
|----|----|
| A  | 1  |
| B  | 2  |
| C  | 3  |

| x1 | x3 |
|----|----|
| A  | T  |
| B  | F  |
| D  | T  |

| x1 | x2 | x3 |
|----|----|----|
| A  | 1  | T  |
| B  | 2  | F  |
| C  | 3  | NA |
| D  | NA | T  |

Mutating Joins

| x1 | x2 | x3 |
|----|----|----|
| A  | 1  | T  |
| B  | 2  | F  |
| C  | 3  | NA |

| x1 | x3 | x2 |
|----|----|----|
| A  | T  | 1  |
| B  | F  | 2  |
| D  | T  | NA |

| x1 | x2 | x3 |
|----|----|----|
| A  | 1  | T  |
| B  | 2  | F  |

| x1 | x2 | x3 |
|----|----|----|
| A  | 1  | T  |
| B  | 2  | F  |
| C  | 3  | NA |
| D  | NA | T  |

Filtering Joins

| x1 | x2 |
|----|----|
| A  | 1  |
| B  | 2  |

| x1 | x2 |
|----|----|
| C  | 3  |

- We often need some different logic to make our joins work. That exists in `dplyr` as well!

y

| x1 | x2 |
|----|----|
| A  | 1  |
| B  | 2  |
| C  | 3  |

+

z

| x1 | x2 |
|----|----|
| B  | 2  |
| C  | 3  |
| D  | 4  |

=

Set Operations

| x1 | x2 |
|----|----|
| B  | 2  |
| C  | 3  |

**`dplyr::intersect(y, z)`**  
Rows that appear in both y and z.

| x1 | x2 |
|----|----|
| A  | 1  |
| B  | 2  |
| C  | 3  |
| D  | 4  |

**`dplyr::union(y, z)`**  
Rows that appear in either or both y and z.

| x1 | x2 |
|----|----|
| A  | 1  |

**`dplyr::setdiff(y, z)`**  
Rows that appear in y but not z.

Binding

| x1 | x2 |
|----|----|
| A  | 1  |
| B  | 2  |
| C  | 3  |
| B  | 2  |
| C  | 3  |
| D  | 4  |


**`dplyr::bind_rows(y, z)`**  
Append z to y as new rows.

| x1 | x2 | x1 | x2 |
|----|----|----|----|
| A  | 1  | B  | 2  |
| B  | 2  | C  | 3  |
| C  | 3  | D  | 4  |

**`dplyr::bind_cols(y, z)`**  
Append z to y as new columns.  
Caution: matches rows by position.

# Joins

- Let's go through our common joins!
- Inner Join: Returns records with matching keys in both tables



# Inner Join

Make our connection and look at the tables

```
library(DBI)
library(dplyr)
con <- dbConnect(RSQLite::SQLite(), "lahman.db")
dbListTables(con)
```

```
## [1] "AllstarFull"      "Appearances"      "AwardsManagers"
## [4] "AwardsPlayers"    "AwardsShareManagers" "AwardsSharePlayers"
## [7] "Batting"          "BattingPost"      "CollegePlaying"
## [10] "Fielding"         "FieldingOF"       "FieldingOFsplit"
## [13] "FieldingPost"     "HallOfFame"       "HomeGames"
## [16] "LahmanData"       "Managers"         "ManagersHalf"
## [19] "Parks"           "People"           "Pitching"
## [22] "PitchingPost"     "Salaries"         "Schools"
## [25] "SeriesPost"       "Teams"            "TeamsFranchises"
## [28] "TeamsHalf"       "battingLabels"    "fieldingLabels"
## [31] "pitchingLabels"
```

# Inner Join

Combine the **Batting** table and the **Pitching** table on common variables

```
inner_join(tbl(con, "Batting") |> filter(yearID == 2000),
           tbl(con, "Pitching") |> filter(yearID == 2000),
           by = join_by(playerID == playerID, stint == stint, teamID == teamID, lgID == lgID)) |>
collect()
```

```
## # A tibble: 677 × 48
##   playerID yearID.x stint teamID lgID   G.x   AB   R.x   H.x   X2B   X3B   HR.x
##   <chr>      <int> <int> <chr> <chr> <int> <int> <int> <int> <int> <int> <int>
## 1 abbotpa...    2000     1 SEA   AL    35     5     1     2     1     0     0
## 2 aceveju...    2000     1 MIL   NL    62     1     1     0     0     0     0
## 3 adamste...    2000     1 LAN   NL    66     2     0     0     0     0     0
## 4 aguilri...    2000     1 CHN   NL    54     0     0     0     0     0     0
## 5 aldresc...    2000     1 PHI   NL    23     0     0     0     0     0     0
## # i 672 more rows
## # i 36 more variables: RBI <int>, SB <int>, CS <int>, BB.x <int>, SO.x <int>,
## #   IBB.x <int>, HBP.x <int>, SH.x <int>, SF.x <int>, GIDP.x <int>,
## #   yearID.y <int>, W <int>, L <int>, G.y <int>, GS <int>, CG <int>, SHO <int>,
## #   SV <int>, IPouts <int>, H.y <int>, ER <int>, HR.y <int>, BB.y <int>,
## #   SO.y <int>, BAOpp <dbl>, ERA <dbl>, IBB.y <int>, WP <int>, HBP.y <int>,
## #   BK <int>, BFP <int>, GF <int>, R.y <int>, SH.y <int>, SF.y <int>, ...
```

# Can Write SQL code instead


- (I'm not a great SQL programmer)

```
tbl(con, sql(
"SELECT p.playerID as pplayerID,
      p.stint as pstint,
      p.teamID as pteamID,
      p.lgID as plgID,
      p.G as pG,
      p.HR as pHR,
      p.BB as pBB,
      p.SO as pSO,
      p.HBP as pHBP,
      p.R as pR,
      p.SF as pSF,
      p.GIDP as pGIDP,
      p.IBB as pIBB,
      p.SH as pSH,
      p.W, p.L, p.GS, p.CG, p.SHO, p.SV, p.IPouts, p.ER, p.BAopp,
      p.ERA, p.WP, p.BK, p.BFP, p.GF,
      b.*
FROM Pitching as p
INNER JOIN Batting as b on ((p.playerID = b.playerID) AND (pstint = b.stint) AND (pteamID = b.teamID) AND (plgID
WHERE b.yearID = 2000 AND p.yearID = 2000"
))
```



# Joins

- Left Join: Returns all records from the 'left' table and any matching records from the 'right' table




# Left Join: Return left table and matching right records

```
left_join(tbl(con, "Batting") |> filter(yearID == 2000),
          tbl(con, "Pitching") |> filter(yearID == 2000),
          by = join_by(playerID == playerID, stint == stint, teamID == teamID, lgID == lgID)) |>
  collect() |>
  select(playerID, ERA, everything())
```

```
## # A tibble: 1,384 × 48
##   playerID  ERA yearID.x stint teamID lgID   G.x   AB   R.x   H.x   X2B   X3B
##   <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <int> <int> <int> <int>
## 1 abbotje... NA      2000     1 CHA    AL     80   215    31    59    15     1
## 2 abbotku... NA      2000     1 NYN    NL     79   157    22    34     7     1
## 3 abbotpa... 4.22    2000     1 SEA    AL     35     5     1     2     1     0
## 4 abreubo... NA      2000     1 PHI    NL    154   576   103   182    42    10
## 5 aceveju... 3.81    2000     1 MIL    NL     62     1     1     0     0     0
## # i 1,379 more rows
## # i 36 more variables: HR.x <int>, RBI <int>, SB <int>, CS <int>, BB.x <int>,
## #   SO.x <int>, IBB.x <int>, HBP.x <int>, SH.x <int>, SF.x <int>, GIDP.x <int>,
## #   yearID.y <int>, W <int>, L <int>, G.y <int>, GS <int>, CG <int>, SHO <int>,
## #   SV <int>, IPouts <int>, H.y <int>, ER <int>, HR.y <int>, BB.y <int>,
## #   SO.y <int>, BAOpp <dbl>, IBB.y <int>, WP <int>, HBP.y <int>, BK <int>,
## #   BFP <int>, GF <int>, R.y <int>, SH.y <int>, SF.y <int>, GIDP.y <int>
```

# Joins

- Right Join: Returns all records from the 'right' table and any matching records from the 'left' table



# Right Join


- Just do a left join and switch the table (or use `right_join()`)

```
right_join(tbl(con, "Batting") |> filter(yearID == 2000),  
           tbl(con, "Pitching") |> filter(yearID == 2000),  
           by = join_by(playerID == playerID, stint == stint, teamID == teamID, lgID == lgID)) |>  
collect() |>  
select(playerID, ERA, everything())
```

```
## # A tibble: 677 × 48  
##   playerID  ERA yearID.x stint teamID lgID   G.x   AB   R.x   H.x   X2B   X3B  
##   <chr>    <dbl>   <int> <int> <chr> <chr> <int> <int> <int> <int> <int> <int>  
## 1 abbotpa... 4.22     2000     1 SEA   AL    35    5     1     2     1     0  
## 2 aceveju... 3.81     2000     1 MIL   NL    62    1     1     0     0     0  
## 3 adamste... 3.52     2000     1 LAN   NL    66    2     0     0     0     0  
## 4 aguilri... 4.91     2000     1 CHN   NL    54    0     0     0     0     0  
## 5 aldresc... 5.75     2000     1 PHI   NL    23    0     0     0     0     0  
## # i 672 more rows  
## # i 36 more variables: HR.x <int>, RBI <int>, SB <int>, CS <int>, BB.x <int>,  
## #   SO.x <int>, IBB.x <int>, HBP.x <int>, SH.x <int>, SF.x <int>, GIDP.x <int>,  
## #   yearID.y <int>, W <int>, L <int>, G.y <int>, GS <int>, CG <int>, SHO <int>,  
## #   SV <int>, IPouts <int>, H.y <int>, ER <int>, HR.y <int>, BB.y <int>,  
## #   SO.y <int>, BAOpp <dbl>, IBB.y <int>, WP <int>, HBP.y <int>, BK <int>,  
## #   BFP <int>, GF <int>, R.y <int>, SH.y <int>, SF.y <int>, GIDP.y <int>
```

# Joins

- Outer Join: Returns all records when there is a match from the 'left' or 'right' table (also called a **full join**)



# Outer Join: Return all matches from both tables

(All players are in the Batting table even if they have no at bats!)

```
full_join(tbl(con, "Batting") |> filter(yearID == 2000),
          tbl(con, "Pitching") |> filter(yearID == 2000),
          by = join_by(playerID == playerID, stint == stint, teamID == teamID, lgID == lgID)) |>
collect()
```

```
## # A tibble: 1,384 × 48
##   playerID yearID.x stint teamID lgID   G.x   AB   R.x   H.x   X2B   X3B   HR.x
##   <chr>      <int> <int> <chr> <chr> <int> <int> <int> <int> <int> <int> <int>
## 1 abbotje...  2000     1  CHA   AL     80  215   31   59   15     1     3
## 2 abbotku...  2000     1  NYN   NL     79  157   22   34     7     1     6
## 3 abbotpa...  2000     1  SEA   AL     35    5     1    2     1     0     0
## 4 abreubo...  2000     1  PHI   NL    154  576  103  182    42    10    25
## 5 aceveju...  2000     1  MIL   NL     62    1     1    0     0     0     0
## # i 1,379 more rows
## # i 36 more variables: RBI <int>, SB <int>, CS <int>, BB.x <int>, SO.x <int>,
## #   IBB.x <int>, HBP.x <int>, SH.x <int>, SF.x <int>, GIDP.x <int>,
## #   yearID.y <int>, W <int>, L <int>, G.y <int>, GS <int>, CG <int>, SHO <int>,
## #   SV <int>, IPouts <int>, H.y <int>, ER <int>, HR.y <int>, BB.y <int>,
## #   SO.y <int>, BAOpp <dbl>, ERA <dbl>, IBB.y <int>, WP <int>, HBP.y <int>,
## #   BK <int>, BFP <int>, GF <int>, R.y <int>, SH.y <int>, SF.y <int>, ...
```

# Other Joins

Those are the major joins covered by `dplyr`. Lots of other joins out there!

- See [here for examples!](#)
  - The right sidebar has more than the standard joins.
- Also ways to do **if then else type logic**, **intersections**, etc. in SQL
- Can do basic **summaries using SQL** as well (including **grouping**), but we'll just use `dplyr` for that!

# Recap

- Joins are combining two tables
- `inner_join` - match records that appear in both tables
- left/right join
- full outer join